Big Data Assignment-5

TEAMMATES:
Rahavi Selvarajan - 1007346445
rahavi.selvarajan@mail.utoronto.ca
Hruday Vishal Kanna Anand - 1006874517
vishal.kanna@mail.utoronto.ca

Part-a

1.

Azure data lake- It is the landing zone for all types of raw data (structured, semi structured, unstructured). It can also be used as our primary storage for all the required data.

Azure Databricks- It is databricks hosted on the azure cloud. It uses spark cluster to transform data in the form of RDDs and data frames.

Azure data factory- is used to transfer data from one service to another or one location to another with the help of linked services.

Azure synapse analytics- is a unified data warehouse platform built on top of azure data lake which takes care of ingestion of data and the analysis of it with the help of SQL or Spark engine.

Azure cosmos db- is a no-SQL data warehouse that is globally distributed and has very low latency.

Ingest data- azure data factory, azure synapse analytics

Data store- azure data lake, Azure cosmos db (noSQL data)

Prepare and transform data- azure databricks, azure synapse analytics

Model and serve data- azure cosmos db, azure synapse analytics

2.

Azure stream analytics has three major steps-

1. Ingest/ input
2. processed
3. Store/ output

In the first step stream analytics ingests data with the help azure event hubs, azure iot hub, and azure blob storage in a continuous manner. This data is then processed; this can be simple SQL queries or

more complex ML services. This processed data might have several types of outputs; it might be in the form of alerts/ actions which will be sent to azure functions, service bus, event hub. It may be used for real time dashboarding; then the data is sent to power BI. The prcessed data may also be stored in a warehouse or normal storage such as data lake or cosmos db.

3.

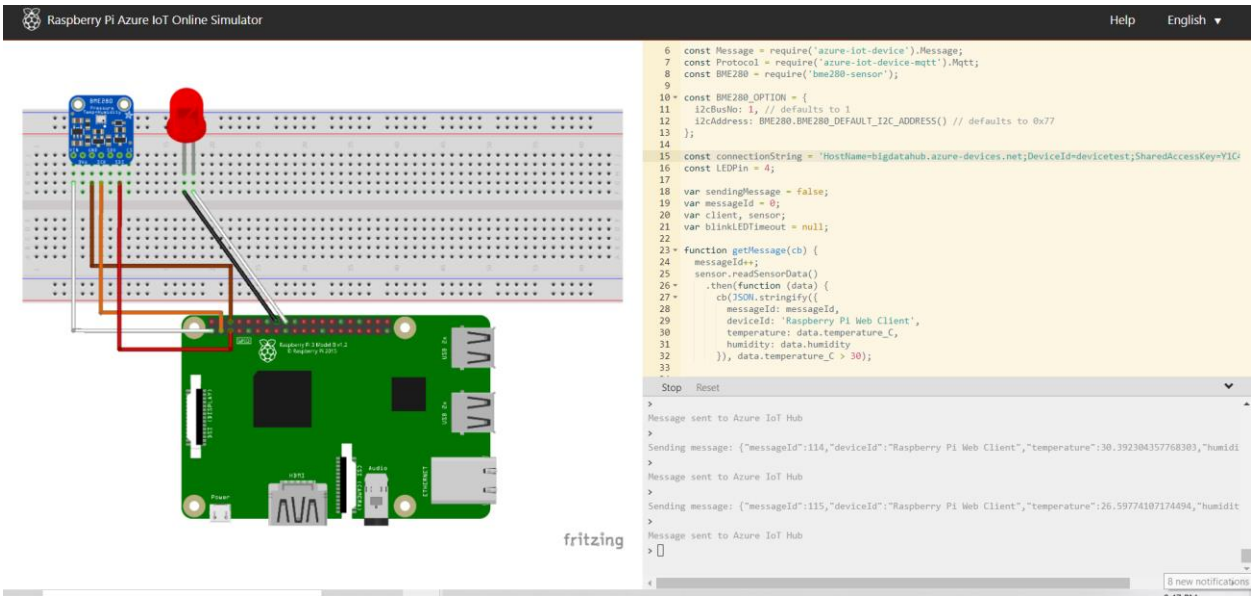

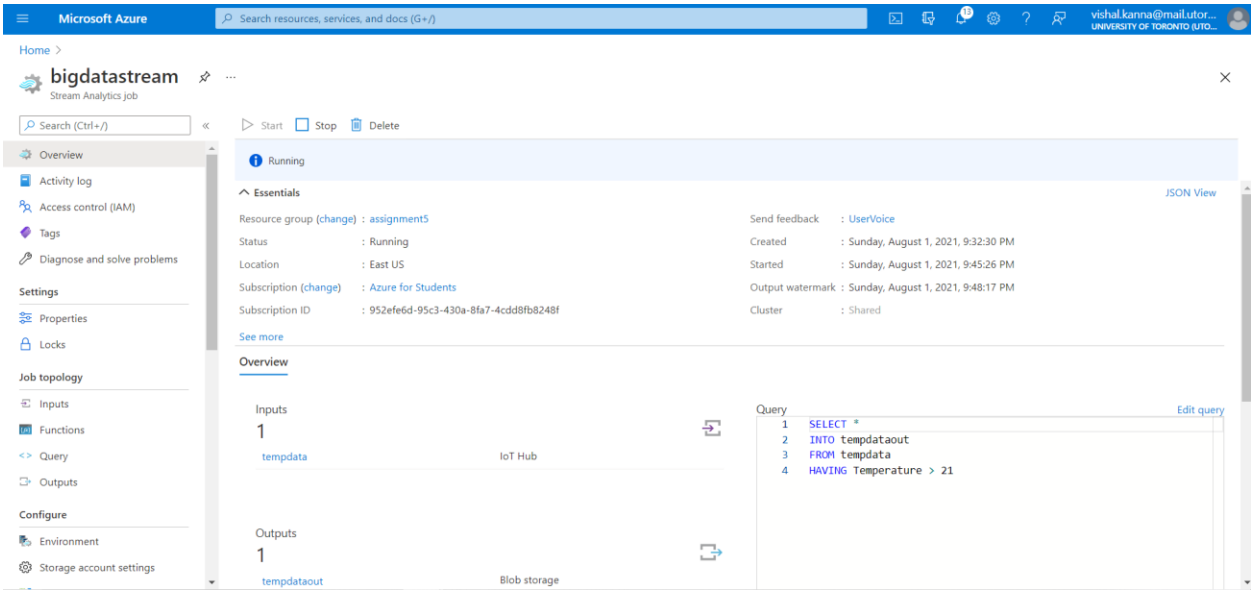

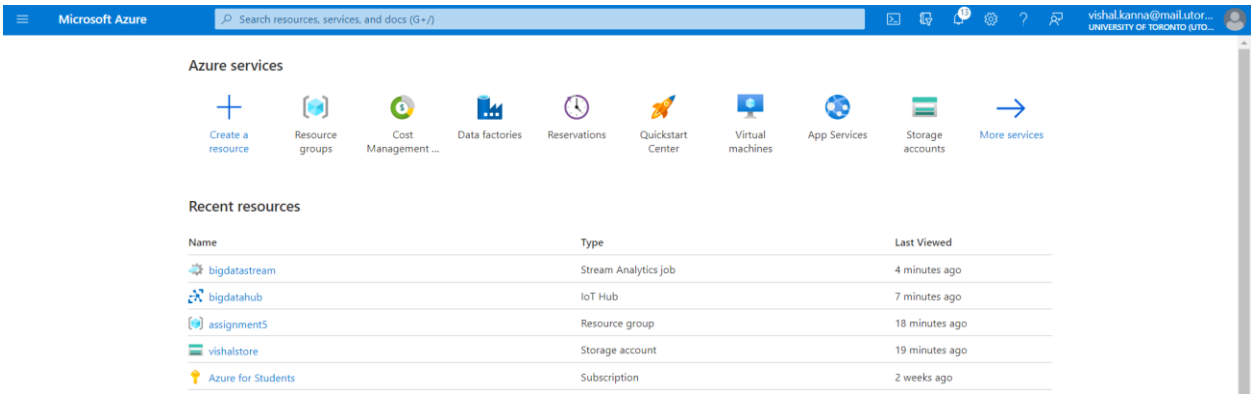
Microsoft Azure
Search resources, services, and docs (G+/)
vishal.kanna@mail.utor...
UNIVERSITY OF TORONTO (UTO...
Azure services
Create a resource
Resource groups
Cost Management ...
Data factories
Reservations
Quickstart Center
Virtual machines
App Services
Storage accounts
More services
Recent resources
Name
Type
Last Viewed
bigdatastream
Stream Analytics job
4 minutes ago
bigdatahub
IoT Hub
7 minutes ago
assignment5
Resource group
18 minutes ago
vishalstore
Storage account
19 minutes ago
Azure for Students
Subscription
2 weeks ago


Search resources, services, and docs (G+/)
vishal.kanna@mail.utor...
UNIVERSITY OF TORONTO (UTO...

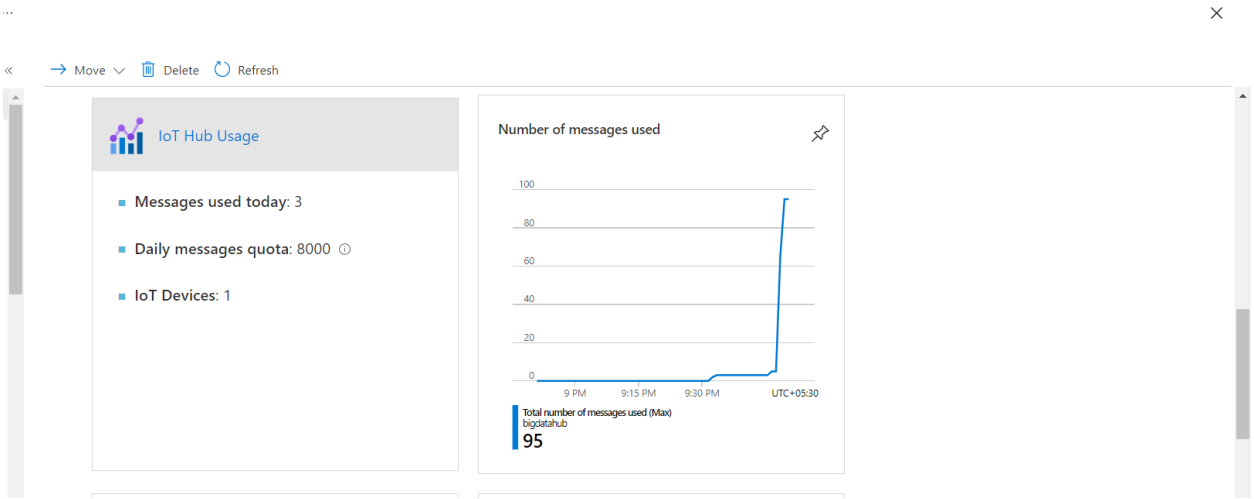
Move
Delete
Refresh
IoT Hub Usage
Messages used today: 3
Daily messages quota: 8000
IoT Devices: 1
Number of messages used
100
80
60
40
20
0
9 PM
9:15 PM
9:30 PM
UTC+05:30
Total number of messages used (Max)
bigdatahub
95

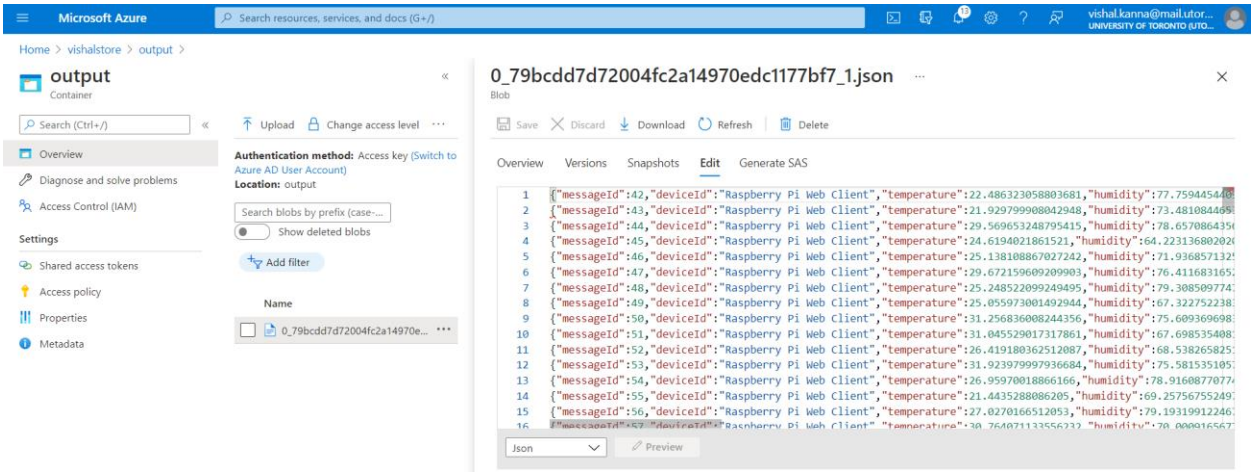
Microsoft Azure
Search resources, services, and docs (G+/)
vishal.kanna@mail.utor...
UNIVERSITY OF TORONTO (UTO...
Home > vishalstore > output >
output
Container
Search (Ctrl+/)
Overview
Diagnose and solve problems
Access Control (IAM)
Settings
Shared access tokens
Access policy
Properties
Metadata
Upload
Change access level
Authentication method: Access key (Switch to Azure AD User Account)
Location: output
Search blobs by prefix (case-...
Show deleted blobs
Add filter
Name
0_79bcdd7d72004fc2a14970e...
0_79bcdd7d72004fc2a14970edc1177bf7_1.json
Blob
Save
Discard
Download
Refresh
Delete
Overview
Versions
Snapshots
Edit
Generate SAS
{"messageId":42,"deviceId":"Raspberry Pi Web Client","temperature":22.486323058803681,"humidity":77.759445403
{"messageId":43,"deviceId":"Raspberry Pi Web Client","temperature":21.929799908042948,"humidity":73.481084465
{"messageId":44,"deviceId":"Raspberry Pi Web Client","temperature":29.569653248795415,"humidity":78.657086435
{"messageId":45,"deviceId":"Raspberry Pi Web Client","temperature":24.6194021861521,"humidity":64.223136802026
{"messageId":46,"deviceId":"Raspberry Pi Web Client","temperature":25.138108867027242,"humidity":71.936857132
{"messageId":47,"deviceId":"Raspberry Pi Web Client","temperature":29.672159609209903,"humidity":76.411683165
{"messageId":48,"deviceId":"Raspberry Pi Web Client","temperature":25.248522099249495,"humidity":79.308509774
{"messageId":49,"deviceId":"Raspberry Pi Web Client","temperature":25.055973001492944,"humidity":67.322752238
{"messageId":50,"deviceId":"Raspberry Pi Web Client","temperature":31.256836008244356,"humidity":75.609369698
{"messageId":51,"deviceId":"Raspberry Pi Web Client","temperature":31.045529017317861,"humidity":67.698535408
{"messageId":52,"deviceId":"Raspberry Pi Web Client","temperature":26.419180362512087,"humidity":68.538265825
{"messageId":53,"deviceId":"Raspberry Pi Web Client","temperature":31.923979997936684,"humidity":75.581535105
{"messageId":54,"deviceId":"Raspberry Pi Web Client","temperature":26.959700188666166,"humidity":78.916087707
{"messageId":55,"deviceId":"Raspberry Pi Web Client","temperature":21.4435288086205,"humidity":69.25756755249
{"messageId":56,"deviceId":"Raspberry Pi Web Client","temperature":27.0270166512053,"humidity":79.193199122461
Json
Preview

```
{"messageId":58,"deviceId":"Raspberry Pi Web Client","temperature":31.887429099256433,"humidity":77.439080212
{"messageId":59,"deviceId":"Raspberry Pi Web Client","temperature":23.733589243233098,"humidity":65.164774819
{"messageId":60,"deviceId":"Raspberry Pi Web Client","temperature":24.527051891653819,"humidity":68.252213826
{"messageId":61,"deviceId":"Raspberry Pi Web Client","temperature":30.5562086339128,"humidity":66.870869552198
{"messageId":62,"deviceId":"Raspberry Pi Web Client","temperature":24.663730123455945,"humidity":76.167345642
{"messageId":63,"deviceId":"Raspberry Pi Web Client","temperature":28.623881163545793,"humidity":72.194303744
{"messageId":64,"deviceId":"Raspberry Pi Web Client","temperature":21.507452893501736,"humidity":66.078687039
{"messageId":65,"deviceId":"Raspberry Pi Web Client","temperature":23.954071724847008,"humidity":66.620100446
{"messageId":68,"deviceId":"Raspberry Pi Web Client","temperature":29.825566797998267,"humidity":60.340044649
{"messageId":69,"deviceId":"Raspberry Pi Web Client","temperature":30.412646579948508,"humidity":72.716832589
{"messageId":70,"deviceId":"Raspberry Pi Web Client","temperature":27.758016975940595,"humidity":70.317070099
{"messageId":71,"deviceId":"Raspberry Pi Web Client","temperature":27.40736970400112,"humidity":78.752502937
{"messageId":72,"deviceId":"Raspberry Pi Web Client","temperature":30.756845959498492,"humidity":69.226669825
{"messageId":73,"deviceId":"Raspberry Pi Web Client","temperature":21.891296329873242,"humidity":78.221896690
{"messageId":74,"deviceId":"Raspberry Pi Web Client","temperature":23.841488923314554,"humidity":66.966216159
```


Json ▾ Preview


```
{"messageId":74,"deviceId":"Raspberry Pi Web Client","temperature":23.841488923314554,"humidity":66.966216159
{"messageId":76,"deviceId":"Raspberry Pi Web Client","temperature":31.779236107648849,"humidity":62.089188818
{"messageId":77,"deviceId":"Raspberry Pi Web Client","temperature":27.825611760378621,"humidity":61.332545484
{"messageId":78,"deviceId":"Raspberry Pi Web Client","temperature":23.595794033027467,"humidity":71.097904818
{"messageId":79,"deviceId":"Raspberry Pi Web Client","temperature":22.686222167849614,"humidity":66.904505974
{"messageId":80,"deviceId":"Raspberry Pi Web Client","temperature":23.563615038479927,"humidity":75.907096602
{"messageId":81,"deviceId":"Raspberry Pi Web Client","temperature":27.915955882678542,"humidity":67.916067745
{"messageId":82,"deviceId":"Raspberry Pi Web Client","temperature":30.727778753966149,"humidity":72.801938404
{"messageId":83,"deviceId":"Raspberry Pi Web Client","temperature":29.565481945985383,"humidity":61.212127547
{"messageId":84,"deviceId":"Raspberry Pi Web Client","temperature":24.996982815569975,"humidity":61.702192290
{"messageId":85,"deviceId":"Raspberry Pi Web Client","temperature":30.125597644991956,"humidity":64.534437428
{"messageId":86,"deviceId":"Raspberry Pi Web Client","temperature":28.196548645180645,"humidity":69.633966779
{"messageId":87,"deviceId":"Raspberry Pi Web Client","temperature":25.22306879709144,"humidity":71.115365741
{"messageId":88,"deviceId":"Raspberry Pi Web Client","temperature":26.559600721636993,"humidity":71.602662458
{"messageId":89,"deviceId":"Raspberry Pi Web Client","temperature":23.497132710919931,"humidity":61.041120087
```



```
{"messageId":88,"deviceId":"Raspberry Pi Web Client","temperature":26.559600721636993,"humidity":71.602662458
{"messageId":89,"deviceId":"Raspberry Pi Web Client","temperature":23.497132710919931,"humidity":61.041120087
{"messageId":90,"deviceId":"Raspberry Pi Web Client","temperature":25.718827026151615,"humidity":67.112418543
{"messageId":91,"deviceId":"Raspberry Pi Web Client","temperature":25.740181131691514,"humidity":61.466858942
{"messageId":92,"deviceId":"Raspberry Pi Web Client","temperature":24.70892037869254,"humidity":77.049711131
{"messageId":93,"deviceId":"Raspberry Pi Web Client","temperature":26.518076218294148,"humidity":79.788552838
{"messageId":94,"deviceId":"Raspberry Pi Web Client","temperature":30.031066108175537,"humidity":73.964708416
{"messageId":95,"deviceId":"Raspberry Pi Web Client","temperature":24.775206107394613,"humidity":77.304864939
```

ASSIGNMENT 5

TEAM MATES:
Rahavi Selvarajan - 1007346445
rahavi.selvarajan@mail.utoronto.ca
Hruday Vishal Kanna Anand - 1006874517
vishal.kanna@mail.utoronto.ca

PART - B

1. **Explain what problem you are going to solve using this dataset. Provide a brief overview of your problem statement.**

   The dataset chosen is **Gait Classification Dataset**. The dataset was created by calculating the walking parameters of 16 different volunteers aged between 20 and 34 years old. A total of 321 attributes are present and out of which only 26 attributes have proper labels.

   Problem Statement: One of the gait attributes is gait variability. Gait variability is the fluctuation of gait measures between steps. Higher variability indicates the risk of falling or frailty or developing a neuro-degenerative disease and is more common in aged people. Lower variability indicates that a person is physically fit. Our target here is gait variability. We are trying to train a model to predict whether a person is prone to the risk of having a neuro degenerative disease or is physically fit with the help of the other gait attributes.

   Cleaning: The dataset has a large number of attributes with no proper labels and these attributes are removed.

   Preprocessing: We have to map the variability feature to 1's and 0's as it will be our target feature for our models.

   Modelling: For the modelling, we chose two algorithms - Logistic Regression and Naive Bayes.

2. **Explain your dataset. Explore your dataset and provide at least 5 meaningful charts/graphs with explanation.**

   1. **Cycle Time Vs Speed**

   Gait cycle time is the time interval between two successive occurrences of one of the repetitive events of the locomotion. Gait Speed is the time one takes to walk a specified distance. From the graph below, we could see that the cycle time remains almost equal for all the volunteers irrespective of the time taken by them to cover a specified distance.

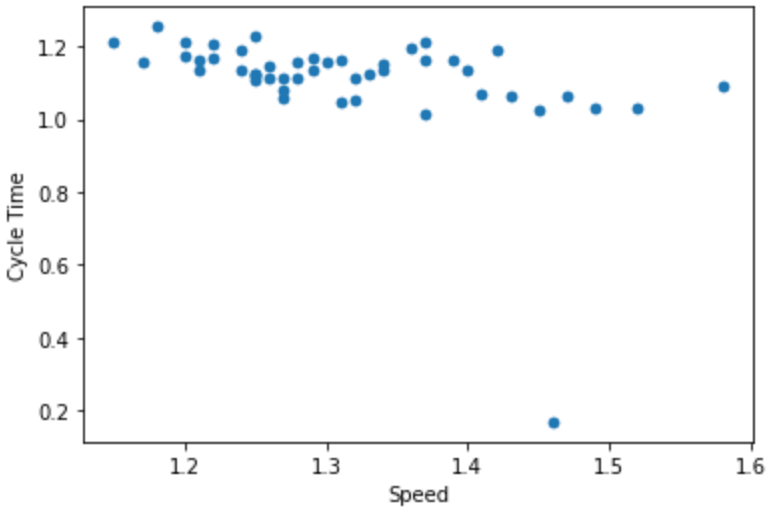

2. Speed Vs Cadence

Gait Cadence is the number of steps or cycles completed by a person in a specified period of time i.e., steps or cycles per minute. From the graph, it is evident that as the speed increases the number of steps taken to complete a specified distance decreases. The graph shows a linearly decreasing behaviour.

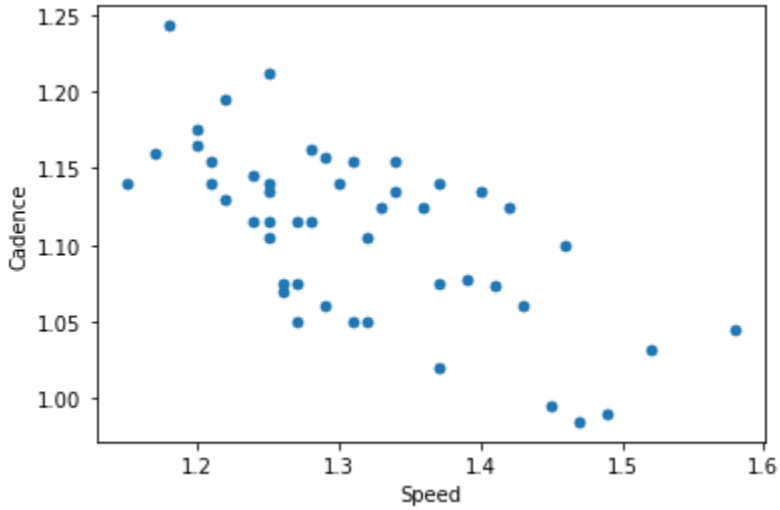

3. Step Angle Grouped by variability

The step angle captures the angular change (ie, rotation) in the sagittal plane of lower limbs during walking. From the graph below, it is evident that people without risks of neurodegenerative diseases have larger step angles (indicated by red) and those with the risks show smaller step angles (indicated by green).

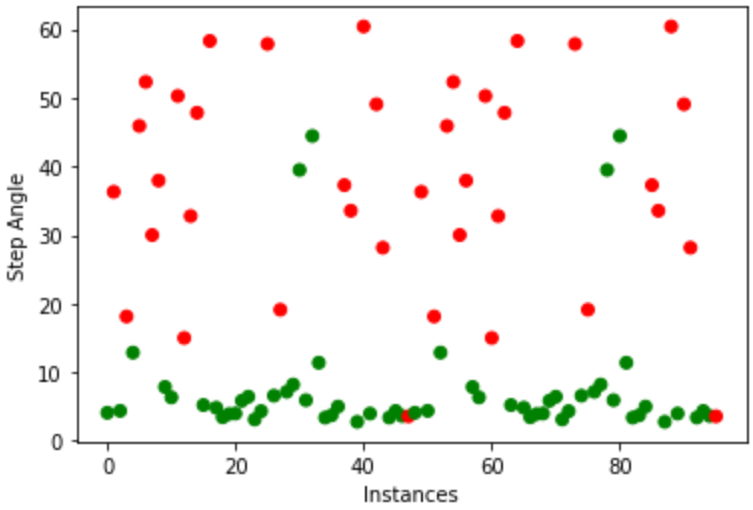

4. Speed Vs Loading - Grouped by Variability

Gait loading is the phase where the body absorbs the impact of the foot and the trunk of the body moves forward to align with the line of gravity. The graph shows that the loading is almost zero in people with the risks of developing the neurodegenerative diseases.

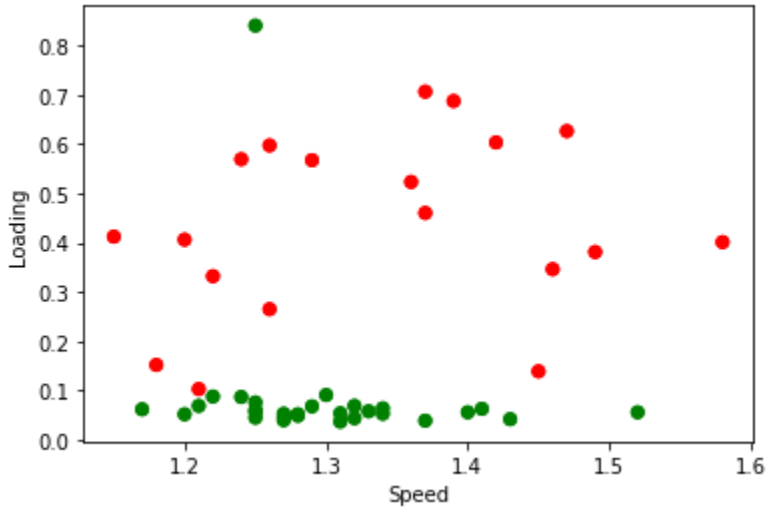

5. **Thrust Vs Loading - Grouped by Variability**

Like the previous graph, irrespective of the thrust, loading remains zero.

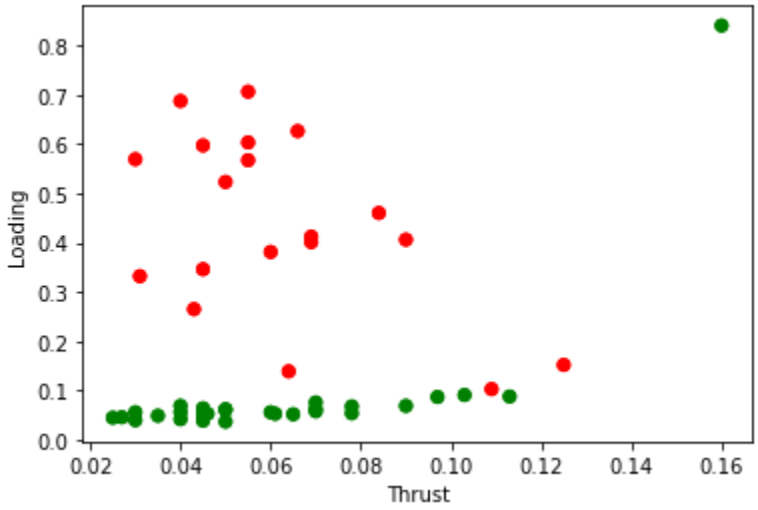

3. **Do data cleaning/pre-processing as required and explain what you have done for your dataset and why?**

Loading dataset into notebook and exploring

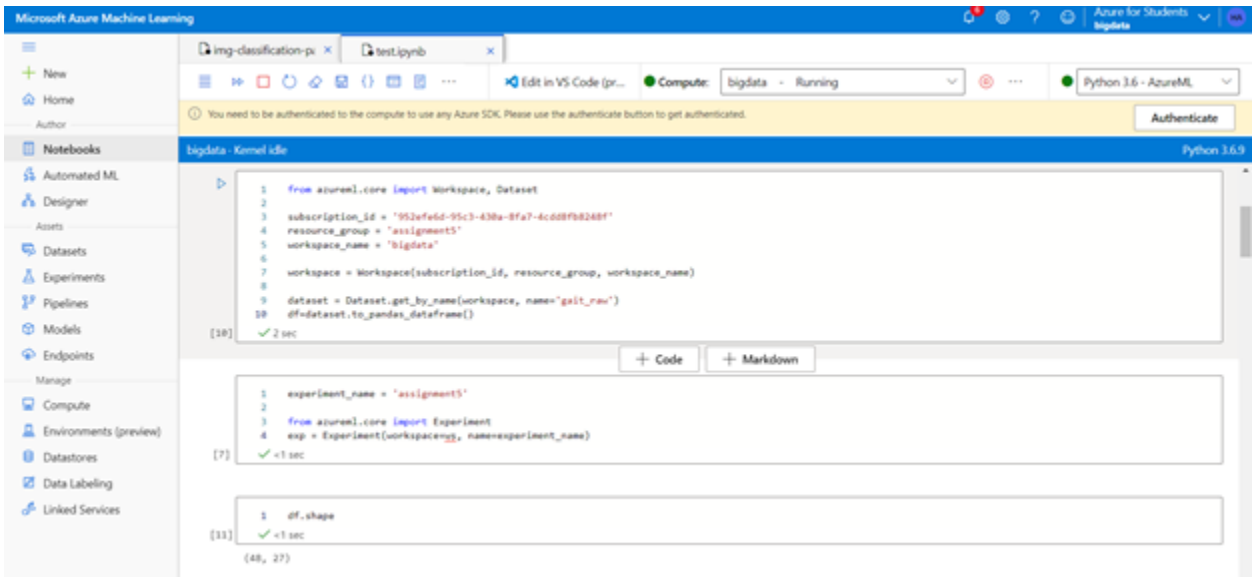

Removing unwanted columns- while we uploaded the file as dataset to azure ml service we removed most of the unlabeled columns but some still persisted. These columns were removed with the code below.

```
[16]  df.drop(["Column172", "Column323"], axis = 1, inplace = True)
      print (df)
      ✓ <1 sec
```

```
    instances  Speed  Variability  Symmetry  Heel Press Time  Cycle Time  \
0           0   1.32         4.15      4.00            1.054       1.054
1           0   1.29         0.00      0.90            1.119       1.137
2           0   1.25         5.06     -3.80            1.109       1.109
3           1   1.21         0.00     -6.30            1.185       1.162
4           1   1.20         4.43     -7.90            1.188       1.172
5           1   1.20         0.00     -8.40            1.235       1.213
6           2   1.37         0.00      1.00            1.239       1.215
7           2   1.46         0.00     -4.50            1.160       0.168
8           2   1.36         0.00      2.80            1.185       1.197
9           3   1.29         5.83     -4.80            1.167       1.170
10          3   1.30         7.84     -5.80            1.161       1.158
11          3   1.58         0.00     -7.40            1.114       1.093
12          4   1.45         0.00      0.90            1.029       1.025
13          4   1.47         0.00      3.50            1.039       1.064
14          4   1.49         0.00      1.80            1.050       1.033
15          5   1.21         6.09      1.30            1.140       1.134
16          5   1.15         0.00     -0.70            1.264       1.211
```

Updating target column(Variability)- to 1 and 0 so that we can proceed to make a valid classification model. All the classification models the target feature needs to be distinct classes for the model to function properly.

```
[17]  df['Variability'] = (df['Variability'] > 0).astype(int)
      df
      ✓ <1 sec
```

	instances	Speed	Variability	Symmetry	Heel Press Time	Cycle Time	Cadence	Posture	Oscillation	Loading	...	Peak Angle Speed,	Max
0	0	1.32	1	4.00	1.054	1.054	1.050	1.060	0.043	0.044	...	1.280	0.99
1	0	1.29	0	0.90	1.119	1.137	1.060	1.065	0.406	0.567	...	5.660	0.98
2	0	1.25	1	-3.80	1.109	1.109	1.105	1.115	0.048	0.056	...	1.265	1.00
3	1	1.21	0	-6.30	1.185	1.162	1.155	1.160	0.215	0.103	...	1.640	1.02
4	1	1.20	1	-7.90	1.188	1.172	1.175	1.177	0.153	0.052	...	1.290	1.01
5	1	1.20	0	-8.40	1.235	1.213	1.165	1.165	0.567	0.406	...	4.345	1.01
6	2	1.37	0	1.00	1.239	1.215	1.140	1.145	0.648	0.460	...	4.655	1.01
7	2	1.46	0	-4.50	1.160	0.168	1.100	1.100	0.348	0.346	...	3.395	0.99

Replicating rows- to create more data for automated ml, as the process requires a minimum of 50 rows to run.

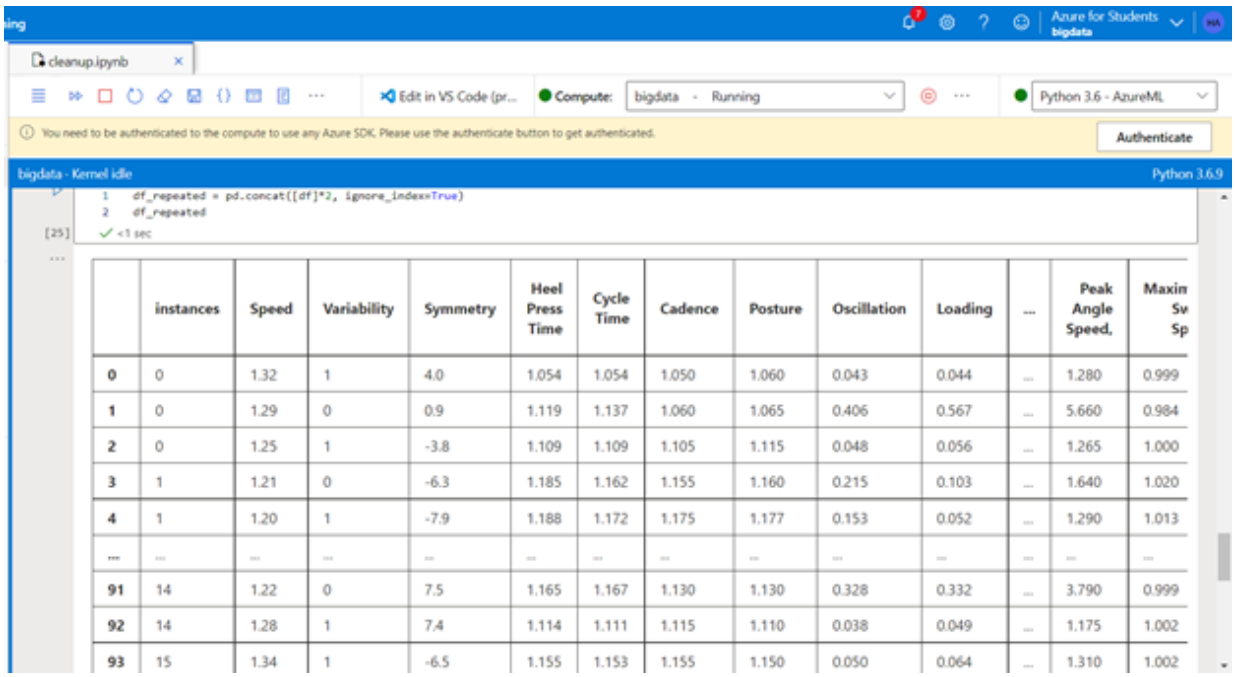

	instances	Speed	Variability	Symmetry	Heel Press Time	Cycle Time	Cadence	Posture	Oscillation	Loading	...	Peak Angle Speed,	Maxin Sv Sp
0	0	1.32	1	4.0	1.054	1.054	1.050	1.060	0.043	0.044	...	1.280	0.999
1	0	1.29	0	0.9	1.119	1.137	1.060	1.065	0.406	0.567	...	5.660	0.984
2	0	1.25	1	-3.8	1.109	1.109	1.105	1.115	0.048	0.056	...	1.265	1.000
3	1	1.21	0	-6.3	1.185	1.162	1.155	1.160	0.215	0.103	...	1.640	1.020
4	1	1.20	1	-7.9	1.188	1.172	1.175	1.177	0.153	0.052	...	1.290	1.013
...	...	...	...	...	...	...	...	...	...	...	...	...	...
91	14	1.22	0	7.5	1.165	1.167	1.130	1.130	0.328	0.332	...	3.790	0.999
92	14	1.28	1	7.4	1.114	1.111	1.115	1.110	0.038	0.049	...	1.175	1.002
93	15	1.34	1	-6.5	1.155	1.153	1.155	1.150	0.050	0.064	...	1.310	1.002

Saving cleaned and processed file-

4. **Implement 2 machine learning models, explain which algorithms you have selected and why. Compare them and show success metrics (Accuracy/RMSE/Confusion Matrix) as per your problem. Explain results.**

Model 1:
Algorithm: Logistic Regression
Regularization Parameter: 2.0
Reason for choosing Logistic Regression: Logistic Regression is the most common algorithm used for binary classification involving numerical values and is known to prevent overfitting.

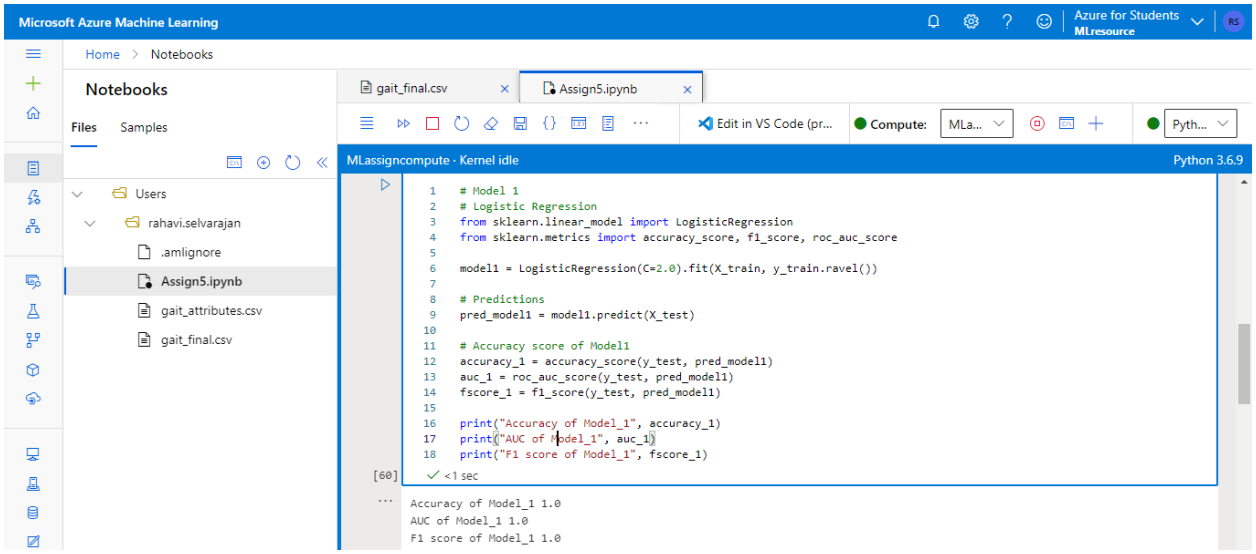

Model 2:

Algorithm: Naive Bayes

Reason for choosing Naive Bayes: It is easy to implement and performs well in test dataset.

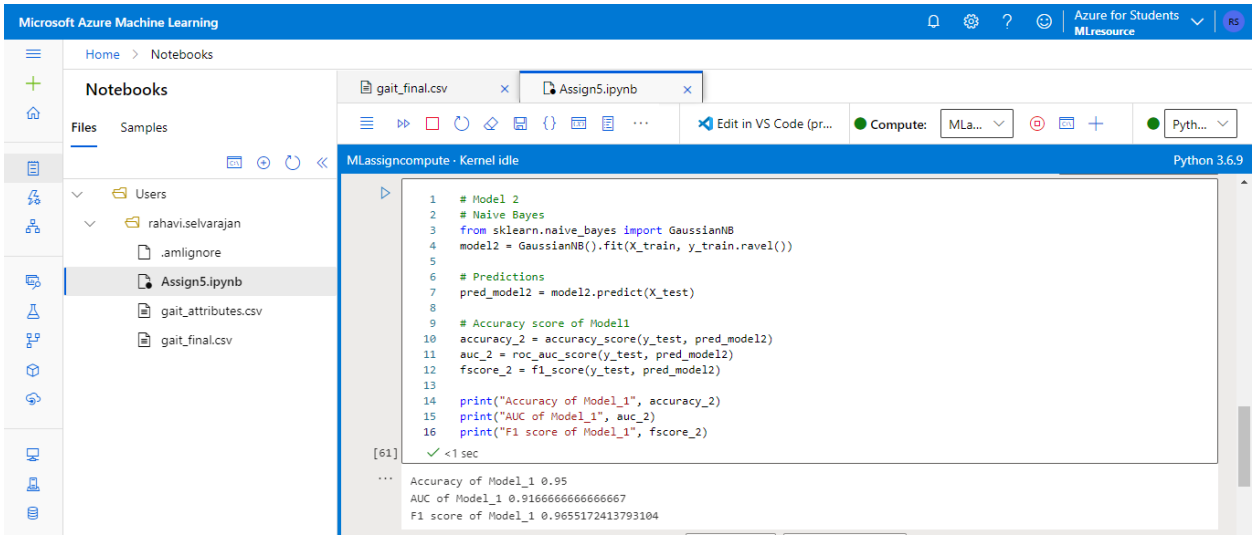

From the above two models, Logistic Regression performs better compared to the Naive Bayes. The Accuracy, F1 score and AUC scores of the first model are better compared to the model 2 i.e, Naive Bayes. F1 score shows the accuracy of the test set and logistic regression shows better performance on the test set.

5. **Use Automated ML for your data set. Explain best model results.**

Automated ml run initiated-

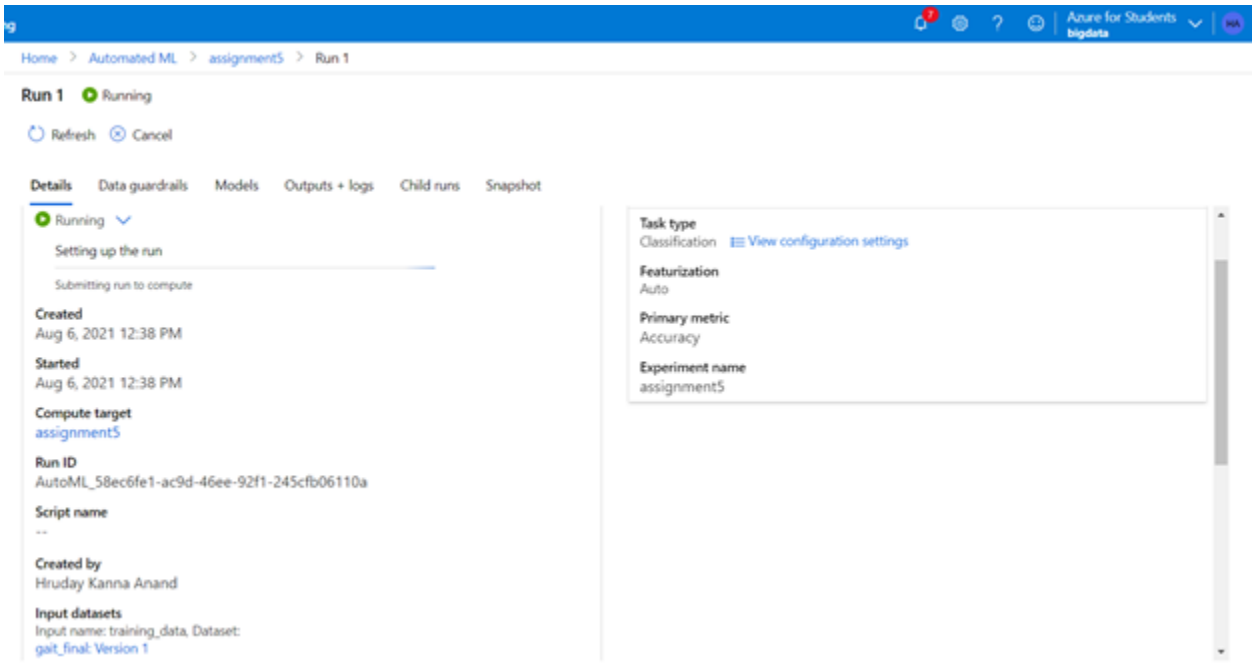

Automated run completed

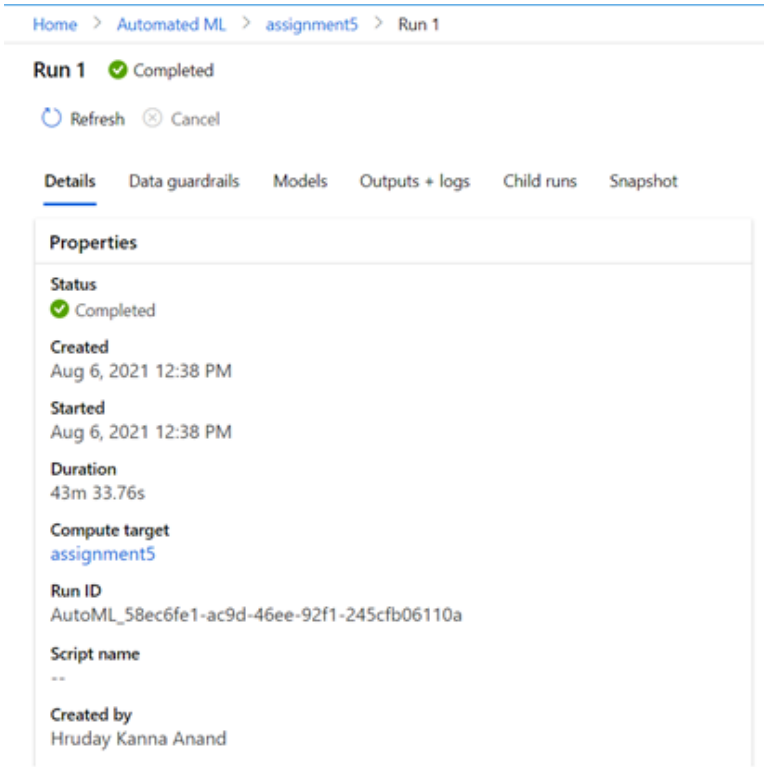

Models-

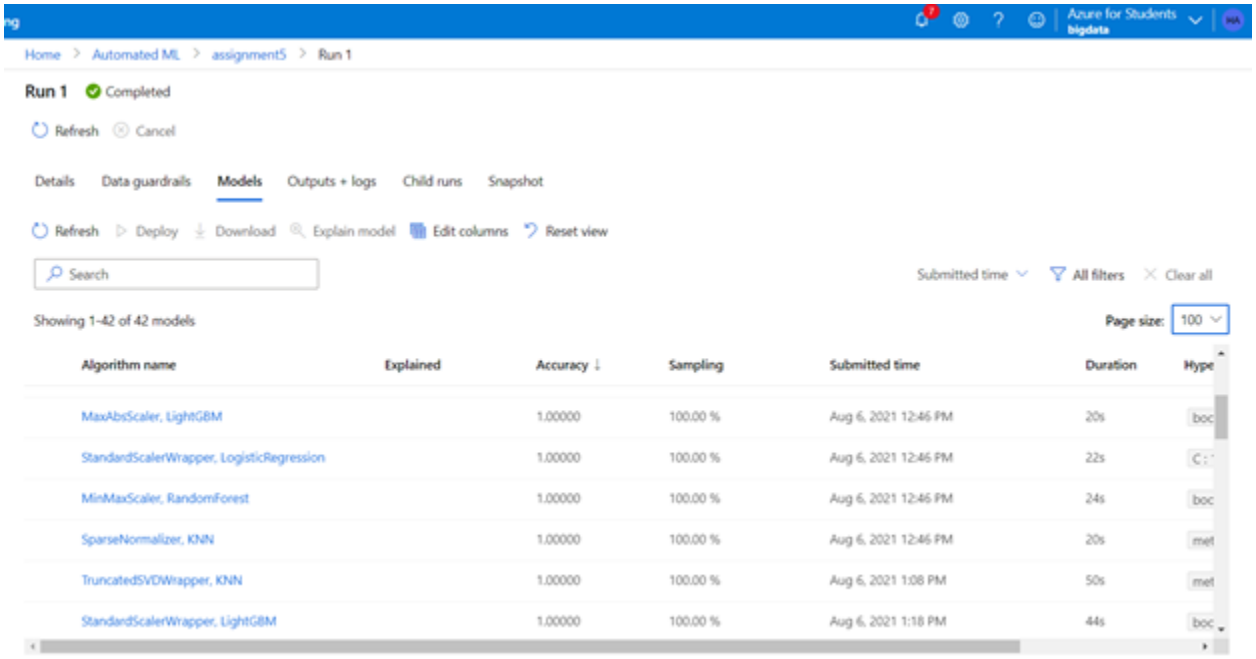

We can see that 42 models have been made and trained and almost all of them have an accuracy of 1.

Best model-

The best model with Light GBM with max abs scaler. Light GBM is a gradient boosting framework that uses tree based learning algorithms. Max abs scaler is used to preprocess our features dividing our features by the maximum value to scale down the features to a common scale.

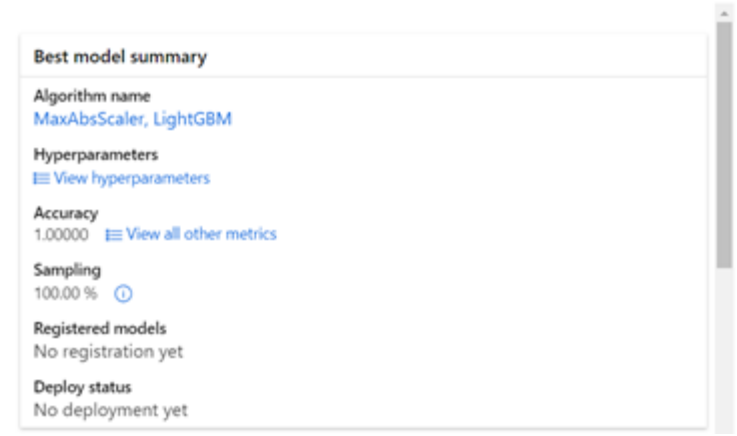

Hyperparameters of the best model-

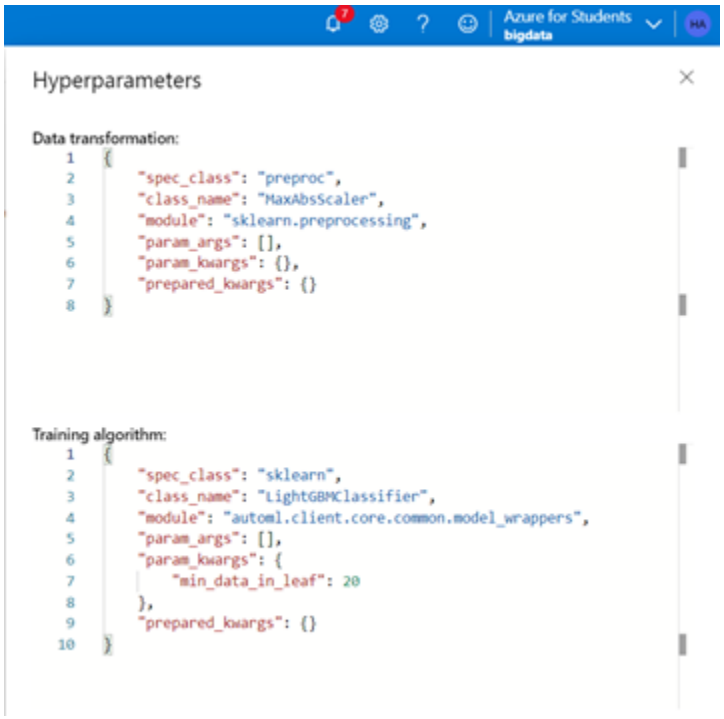

Model explanation-

We can see the model has an accuracy, precision and recall 1.

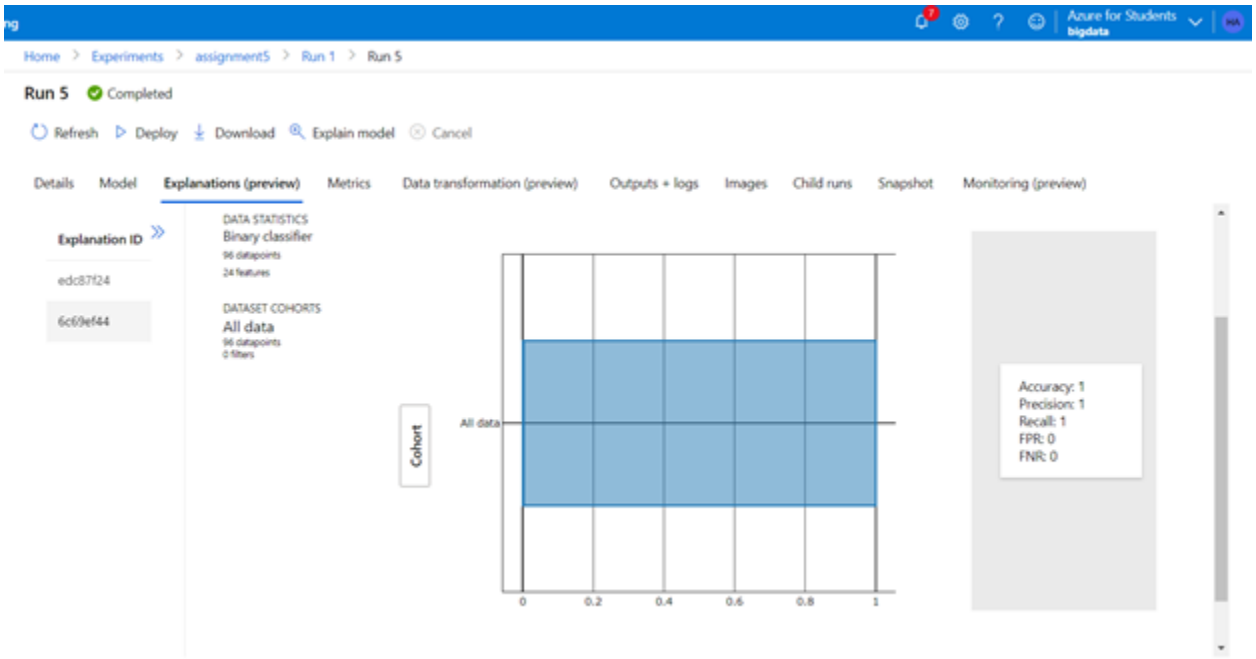

From this graph we can see our model relies largely on the peak angle speed feature to come up with an accurate output.